

1.0-jo-initial-data-exploration

April 21, 2022

1 Gender Pay Gap

Team members: * Jordan Farrell * Julio Oliveira * Ashwath Ramesh * Satoshi Taniguchi * Junjun Tao * William Teodecki

```
[ ]: import numpy as np
import pandas as pd

import shap
from fitter import Fitter
import statsmodels.api as sm
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression

from sklearn.metrics import r2_score, mean_squared_error
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split

import seaborn as sns
import matplotlib.pyplot as plt
```

```
/home/julio/.local/share/virtualenvs/gender-pay-gap-VcKyyecI/lib/python3.8/site-
packages/tqdm/auto.py:22: TqdmWarning: IProgress not found. Please update
jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user\_install.html
from .autonotebook import tqdm as notebook_tqdm
```

```
[ ]: columns = [
    'sex', 'age', 'annhrs', 'annlabinc', 'white', 'black', 'hisp', 'degree',
    'yrsexp', 'yrsftexp', 'yrsptexp', 'hrwage', 'northeast', 'northcentral',
    'south', 'west', 'manager', 'business', 'financialop', 'computer',
    'architect', 'scientist', 'socialworker', 'postseceduc', 'legaleduc',
    'artist', 'lawyerphysician', 'healthcare', 'healthsupport', 'protective',
    'foodcare', 'building', 'sales', 'officeadmin', 'constructextractinstall',
    'production', 'transport', 'farminc', 'labincbus'
```

```
]

psid = pd.read_csv('../data/external/PanelStudyIncomeDynamics.csv',
                    usecols=columns)
```

```
[ ]: psid.head()
```

```
[ ]:
   sex  farminc  age  annhrs  annlabinc  white  black  hisp  degree  \
0    1      0.0   34   1600   10000.0     1     0     0     1.0
1    1      0.0   32    520    9095.0     0     1     0     0.0
2    1      0.0   64   2550   45200.0     0     1     0     0.0
3    1      0.0   50   3072   25000.0     1     0     0     0.0
4    1      0.0   26   2100   24500.0     1     0     0     0.0

   labincbus  ...  healthcare  healthsupport  protective  foodcare  building  \
0           0  ...           0              0           0           0           0
1           0  ...           0              0           0           0           0
2           0  ...           0              0           0           0           0
3           0  ...           0              0           0           0           0
4           0  ...           0              0           0           0           0

   sales  officeadmin  constructextractinstall  production  transport
0       1           0              0           0           0
1       0           0              0           1           0
2       0           0              0           1           0
3       0           0              0           0           1
4       0           0              0           1           0
```

```
[5 rows x 39 columns]
```

1.0.1 Column Encoding

```
[ ]: def get_category(row, columns):
      for col in columns:
          if row[col] == 1:
              return col

race_cols = ['white', 'black', 'hisp']
region_cols = ['northeast', 'northcentral', 'south', 'west']
job_cols = [
    'manager', 'business', 'financialop', 'computer', 'architect', 'scientist',
    'socialworker', 'postseceduc', 'legaleduc', 'artist', 'lawyerphysician',
    'healthcare', 'healthsupport', 'protective', 'foodcare', 'building',
    'sales', 'officeadmin', 'constructextractinstall', 'production',
    'transport'
]
```

```

psid['race'] = psid.apply(lambda x: get_category(x, race_cols), axis=1)
psid.drop(race_cols, axis=1, inplace=True)

psid['region'] = psid.apply(lambda x: get_category(x, region_cols), axis=1)
psid.drop(region_cols, axis=1, inplace=True)

psid['job'] = psid.apply(lambda x: get_category(x, job_cols), axis=1)
psid.drop(job_cols, axis=1, inplace=True)

psid['sex'] = psid.apply(lambda x: 'male' if x['sex'] == 1 else 'female',
                        axis=1)

degree_map = {0: 'no_college', 1: 'bachelors', 2: 'advanced_degree'}
psid['degree'] = psid.degree.map(degree_map)

```

```
[ ]: psid.shape
```

```
[ ]: (33398, 14)
```

```
[ ]: psid.head()
```

```

[ ]:
   sex  farminc  age  annhrs  annlabinc  degree  labincbus  yrsexp \
0  male      0.0   34   1600   10000.0  bachelors         0    12.0
1  male      0.0   32    520    9095.0 no_college         0    14.0
2  male      0.0   64   2550   45200.0 no_college         0    39.0
3  male      0.0   50   3072   25000.0 no_college         0    30.0
4  male      0.0   26   2100   24500.0 no_college         0     8.0

   yrsfexp  yrsptexp   hrwage  race  region  job
0     12.0      0.0  6.250000  white  northeast  sales
1     11.0      3.0  17.490385  black  northcentral  production
2     38.0      1.0  17.725491  black  northeast  production
3     30.0      0.0   8.138021  white  northcentral  transport
4      8.0      0.0  11.666667  white      south  production

```

1.0.2 Income Columns

```

[ ]: income_cols = ['farminc', 'annlabinc', 'labincbus', 'annhrs', 'hrwage']
psid[income_cols].describe()

```

```

[ ]:
   count  farminc  annlabinc  labincbus  annhrs  hrwage
count  33398.000000  3.339800e+04  33398.000000  33398.000000  33398.000000
mean    104.879544  3.708689e+04   168.051859   1990.103449   18.418722
std     2662.001098  4.156487e+04   2257.977695   623.592732   19.462814
min     -5000.000000  3.000000e+01    0.000000    10.000000    0.891473
25%         0.000000  1.600000e+04    0.000000   1767.000000    8.823529

```

50%	0.000000	2.900000e+04	0.000000	2000.000000	14.423077
75%	0.000000	4.600000e+04	0.000000	2277.000000	22.373541
max	200000.000000	1.500000e+06	99999.000000	5840.000000	1000.000000

Looking into anomalous values in income we were able to find the farm incomes can be negative.

```
[ ]: psid[psid['farminc'] != 0][income_cols].describe()
```

```
[ ]:
count      farminc      annlabinc      labincbus      annhrs      hrwage
mean    29189.725000    24491.858333    114.133333    1860.741667    12.721140
std     33654.453819    18135.461596    933.531183    668.718635     7.682096
min     -5000.000000     400.000000     0.000000    129.000000     2.334267
25%      7000.000000    10748.750000     0.000000    1470.000000     7.179206
50%     20000.000000    21000.000000     0.000000    2000.000000    10.971956
75%     40250.000000    33325.000000     0.000000    2173.250000    16.811773
max     200000.000000   101400.000000   10000.000000   3491.000000    41.666668
```

100% of the dataset has annual labor income greater than 0.

```
[ ]: psid[psid.annlabinc > 0].shape[0] / psid.shape[0]
```

```
[ ]: 1.0
```

There is no considerable amount of missing values for the main attributes.

```
[ ]: psid.isnull().sum()
```

```
[ ]: sex          0
     farminc      0
     age          0
     annhrs       0
     annlabinc    0
     degree      40
     labincbus    0
     yrsexp       0
     yrsftexp     0
     yrsptexp     0
     hrwage       0
     race        417
     region       3
     job         189
     dtype: int64
```

1.0.3 Classes distribution

See dataset description for more details: <https://github.com/jcalvesoliveira/gender-pay-gap/blob/master/docs/datasets/PanelStudyIncomeDynamics.names>


Distribution of labor income

```
[ ]: f = Fitter(psid.annlabinc,
              distributions=['lognorm',
                           "norm", "expon", "chi2"])

f.fit()
f.summary()
```

```
[ ]:      sumsquare_error      aic      bic  kl_div
chi2      1.402789e-11    8022.927868 -1.182466e+06    inf
expon      8.829477e-11    6155.790778 -1.121036e+06    inf
norm       2.396454e-10   42587.165871 -1.087689e+06    inf
lognorm    7.046555e-10   3425.914448 -1.051657e+06    inf
```

WARNING:matplotlib.font_manager.findfont: Font family ['sans-serif'] not found.
Falling back to DejaVu Sans.



The labor income distribution is skewed to the right, meaning that a small amount of people will make a lot of money. In order to focus our analysis in the overall population we filter out the 10% highest incomes by selecting the 90th percentile.

```
[ ]: q_99 = psid.annlabinc.quantile(0.99)
      q_99
```

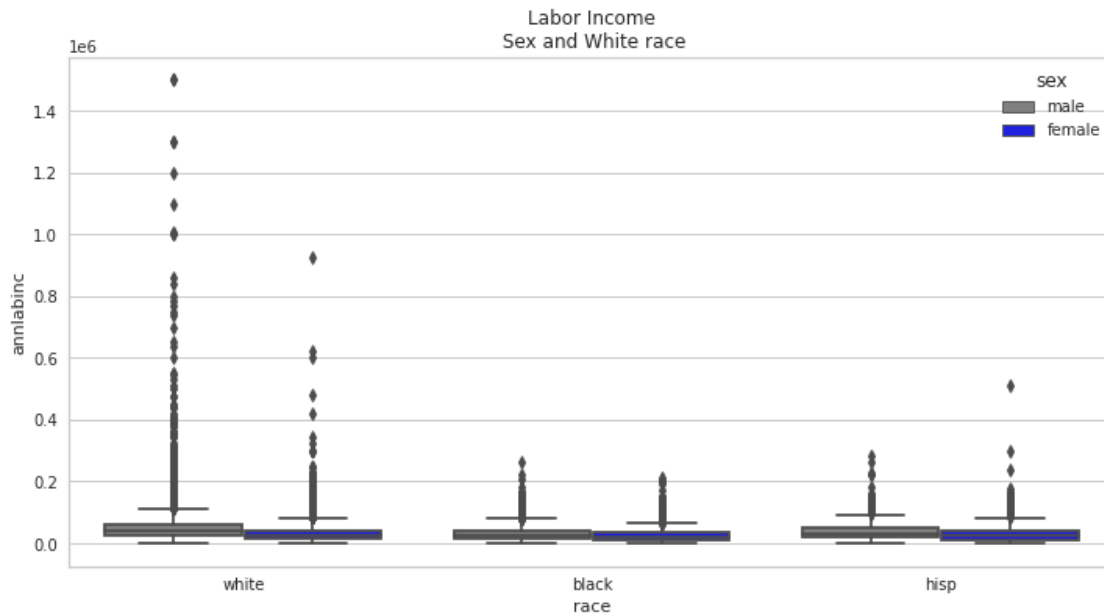
```
[ ]: 171000.0
```

1.0.4 What categorical features may have an effect on income?

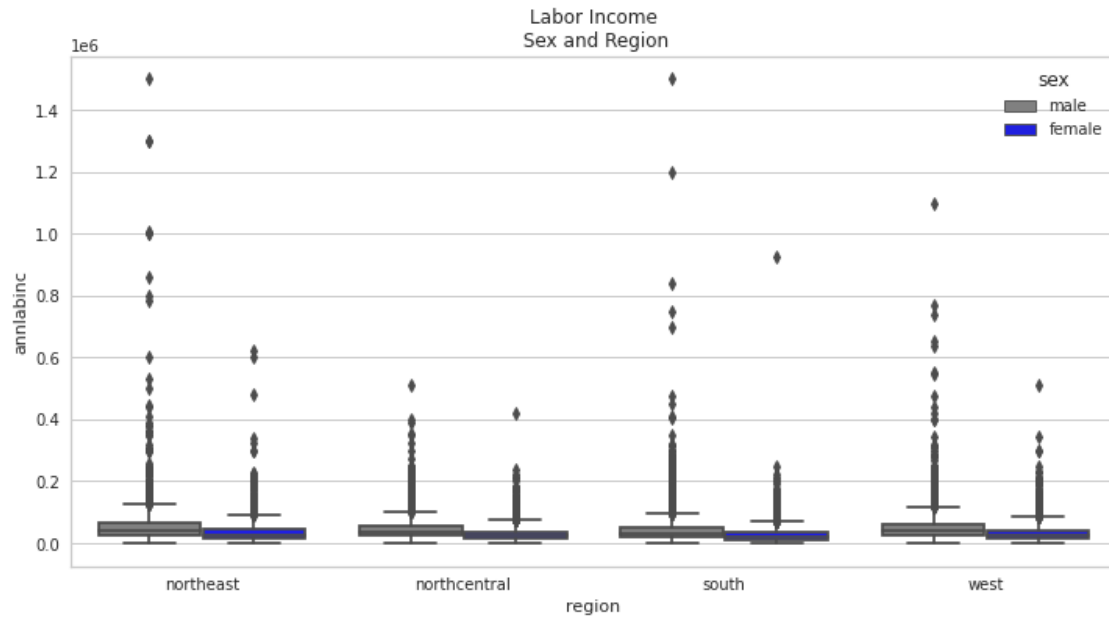
```
[ ]: fig, ax = plt.subplots(figsize=(12, 6))
ax.set(title="Labor Income\n Sex and White race")
ax = sns.boxplot(x="race", y="annlabinc", hue="sex",
                data=psid,
                dodge=True, ax = ax, palette=['gray', 'blue'])
```

WARNING:matplotlib.font_manager.findfont: Font family ['sans-serif'] not found.
Falling back to DejaVu Sans.


WARNING:matplotlib.font_manager.findfont: Font family ['sans-serif'] not found.
Falling back to DejaVu Sans.



```
[ ]: fig, ax = plt.subplots(figsize=(12, 6))
ax.set(title="Labor Income\n Sex and Region")
ax = sns.boxplot(x="region", y="annlabinc", hue="sex",
                data=psid,
                dodge=True, ax = ax, palette=['gray', 'blue'])
```



```
[ ]: fig, ax = plt.subplots(figsize=(12, 6))
ax.set(title="Labor Income\n Degree level")
ax = sns.boxplot(x="degree", y="annlabinc", hue="sex",
                data=psid,
                dodge=True, ax = ax, palette=['gray', 'blue'])
```




1.1 Data Cleaning

```
[ ]: psid = psid[psid.annlabinc < q_99]
psid.degree.fillna('no_college', inplace=True)
psid.job.fillna('na', inplace=True)
psid.race.fillna('na', inplace=True)
psid = psid[psid['region'].notna()]
```

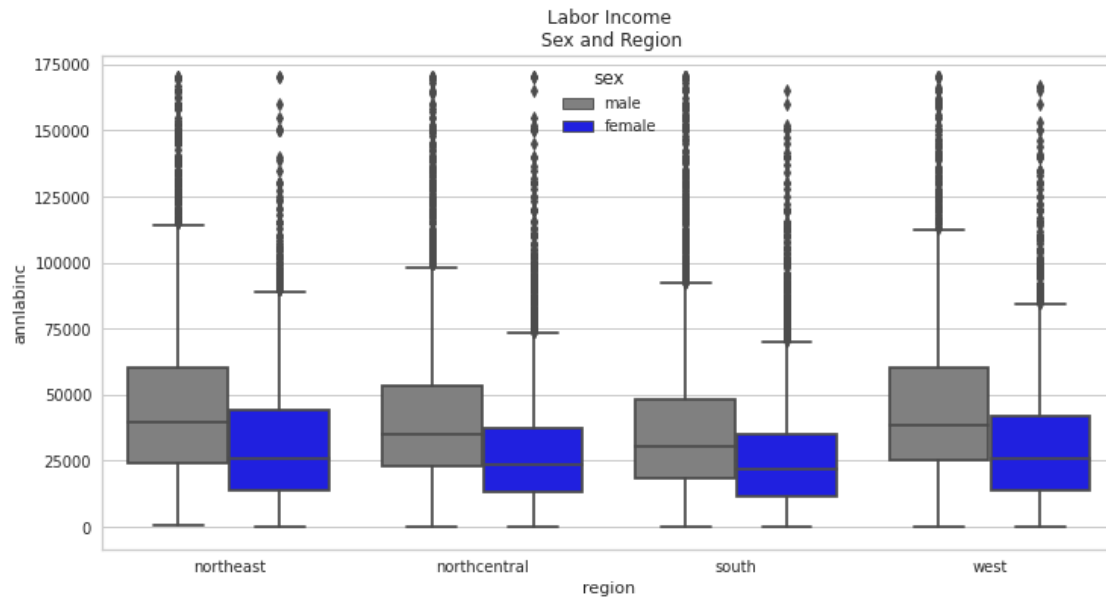
```
[ ]: psid.to_csv('../data/processed/psid.csv')
```

1.1.1 What categorical features may have an effect on income?


```
[ ]: fig, ax = plt.subplots(figsize=(12, 6))
ax.set(title="Labor Income\n Sex and White race")
ax = sns.boxplot(x="race", y="annlabinc", hue="sex",
                 data=psid,
                 dodge=True, ax = ax, palette=['gray', 'blue'])
```



```
[ ]: fig, ax = plt.subplots(figsize=(12, 6))
ax.set(title="Labor Income\n Sex and Region")
ax = sns.boxplot(x="region", y="annlabinc", hue="sex",
                 data=psid,
                 dodge=True, ax = ax, palette=['gray', 'blue'])
```

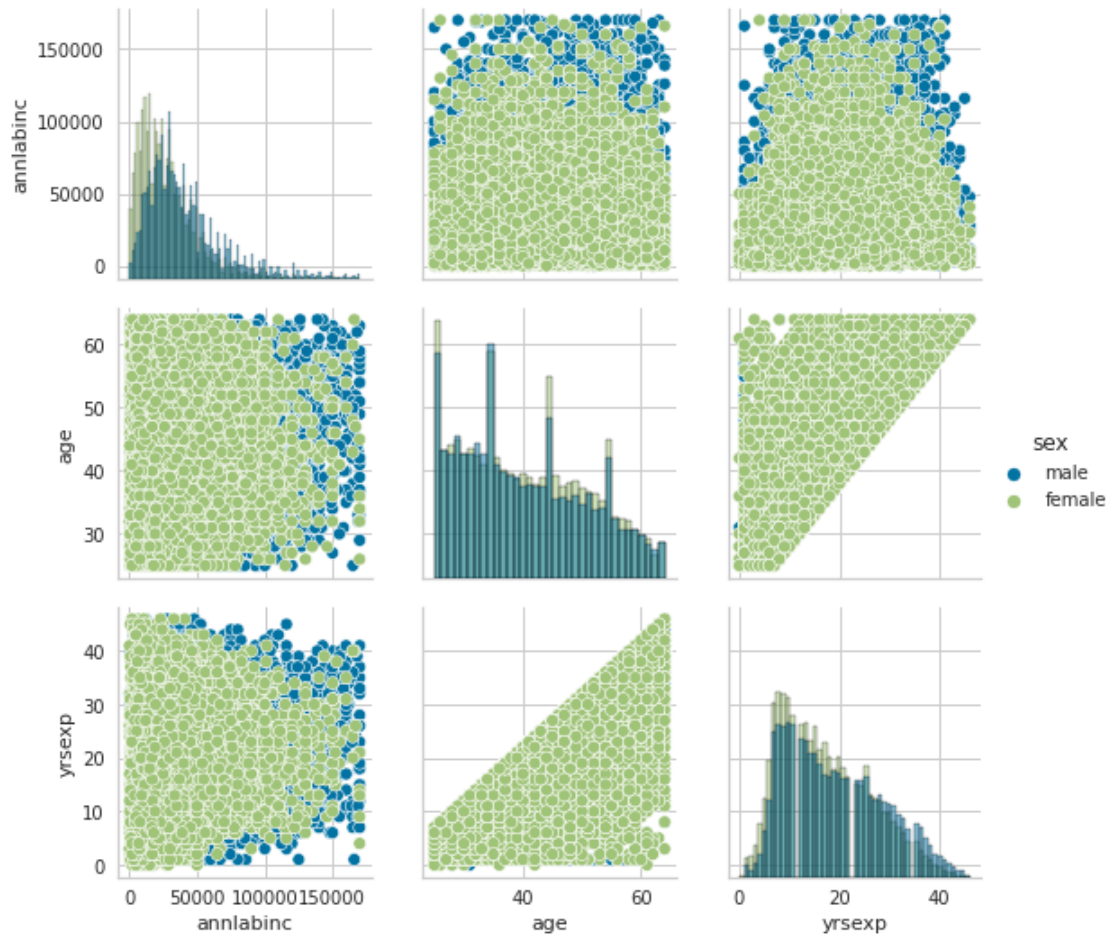
```
[ ]: fig, ax = plt.subplots(figsize=(12, 6))
ax.set(title="Labor Income\n Degree level")
ax = sns.boxplot(x="degree", y="annlabinc", hue="sex",
                data=psid,
                dodge=True, ax = ax, palette=['gray', 'blue'])
```




Looking into the plot below, we can observe that different from what we would expect we could not identify a clear relationship between age and years of experience with labor income.

```
[ ]: g = sns.PairGrid(psid[['annlabinc', 'age', 'yrsexp']], hue="sex")
g.map_diag(sns.histplot)
g.map_offdiag(sns.scatterplot)
g.add_legend()
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x7f99e2f51be0>
```



```
[ ]: plt.figure(figsize=(16, 6))
heatmap = sns.heatmap(psid[['annlabinc', 'age', 'yrsexp']].corr(), vmin=-1,
    ↪vmax=1, annot=True)
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':12}, pad=12);
```



2 Data Transformation

```
[ ]: psid.drop(['annlabinc', 'yrsftexp', 'yrsptexp', 'farminc', 'labincbus'], axis=1,
               ↪inplace=True)
```

```
[ ]: psid.age.describe()
```

```
[ ]: count    33060.000000
      mean      40.411827
      std      10.397327
      min      25.000000
      25%      31.000000
      50%      39.000000
      75%      49.000000
      max      64.000000
      Name: age, dtype: float64
```

```
[ ]: psid.annhrs.describe()
```

```
[ ]: count    33060.000000
      mean    1985.416243
      std     623.083807
      min      10.000000
      25%    1764.000000
      50%    2000.000000
      75%    2268.000000
      max    5840.000000
      Name: annhrs, dtype: float64
```

```
[ ]: psid.yrsexp.describe()
```

```
[ ]: count    33060.000000
      mean      18.061162
      std       9.383590
      min       0.000000
      25%      10.000000
      50%      17.000000
      75%      25.000000
      max      46.000000
      Name: yrsexp, dtype: float64
```

```
[ ]: psid['annhrs'] = pd.cut(psid['annhrs'], bins=[0,1764, 2000,
      ↪2268,6000],labels=['0-1764','1764-2000','2000-2268','2268-6000'])
psid['age'] = pd.cut(psid['age'], bins=[0,31, 39,
      ↪49,64],labels=['0-31','31-39','39-49','49-64'])
psid['yrsexp'] = pd.cut(psid['yrsexp'], bins=[0,10, 17, 25,46],
      ↪labels=['0-10','10-17','17-25','25-46'])
```

```
[ ]: X = psid.drop(['hrwage'], axis=1)
      y = psid.hrwage.copy()
      X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
      ↪random_state=42)
      print(f"X_train shape: {X_train.shape}, X_test shape: {X_test.shape},
      ↪proportion: {X_train.shape[0] / (X_train.shape[0] + X_test.shape[0])}")
```

```
X_train shape: (26448, 8), X_test shape: (6612, 8), proportion: 0.8
```

```
[ ]: cat_cols = ['sex','degree','race', 'job', 'region','age','annhrs','yrsexp']
```

```
[ ]: encoding_transf = ColumnTransformer(remainder='passthrough',
      transformers=[('ohe', OneHotEncoder(drop='first',
      ↪sparse=False), cat_cols)],
      verbose_feature_names_out=True)
norm_transf = ColumnTransformer(remainder='passthrough',
      transformers=[('norm', MinMaxScaler())])

pipeline = Pipeline([
    ('encoding',encoding_transf),
    ('normalization', MinMaxScaler())
])
```

2.0.1 Modeling

```
[ ]: X_train = pipeline.fit_transform(X_train,y_train)
X_test = pipeline.transform(X_test)
lr = LinearRegression()
lr.fit(X_train, y_train)
```

```
[ ]: LinearRegression()
```

```
[ ]: feature_names=list(pipeline.get_feature_names_out())
```

```
[ ]: predicted = lr.predict(X_test)
predicted
```

```
[ ]: array([ 6.86714946,  9.76743895,  9.52532006, ...,  8.36140376,
          31.97124922, 16.48811779])
```

```
[ ]: print(f"Hourly wage mean:{y_test.mean():.2f}, std:{y_test.std():.2f}, range:
      ↳{y_test.min():.2f}-{y_test.max():.2f}")
print(f"R2 score: {r2_score(y_test, predicted)}")
print(f"MSE: {mean_squared_error(y_test, predicted)}")
print(f"RMSE: {np.sqrt(mean_squared_error(y_test, predicted))}")
```

Hourly wage mean:17.18, std:12.40, range:0.89-166.56
R2 score: 0.2645586985709846
MSE: 113.07469607048867
RMSE: 10.633658639926743

```
[ ]: X2 = sm.add_constant(X_train)
ols = sm.OLS(y_train, X2).fit()
print(ols.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          hrwage      R-squared:                0.209
Model:                  OLS        Adj. R-squared:            0.208
Method:                 Least Squares    F-statistic:          174.3
Date:                  Thu, 21 Apr 2022    Prob (F-statistic):      0.00
Time:                  18:58:08      Log-Likelihood:        -1.0523e+05
No. Observations:      26448      AIC:                  2.105e+05
Df Residuals:          26407      BIC:                  2.109e+05
Df Model:               40
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	25.8702	0.701	36.892	0.000	24.496	27.245
x1	3.7306	0.199	18.767	0.000	3.341	4.120

x2	-4.5190	0.369	-12.248	0.000	-5.242	-3.796
x3	-11.7668	0.364	-32.314	0.000	-12.481	-11.053
x4	1.1109	0.489	2.274	0.023	0.153	2.069
x5	-0.4050	0.746	-0.543	0.587	-1.867	1.057
x6	0.9703	0.195	4.987	0.000	0.589	1.352
x7	-3.0902	0.907	-3.406	0.001	-4.868	-1.312
x8	-9.8649	0.691	-14.274	0.000	-11.220	-8.510
x9	0.3798	0.772	0.492	0.623	-1.133	1.892
x10	2.8755	0.763	3.770	0.000	1.381	4.370
x11	-4.1219	0.603	-6.840	0.000	-5.303	-2.941
x12	0.1491	0.811	0.184	0.854	-1.441	1.740
x13	-8.4714	0.641	-13.222	0.000	-9.727	-7.216
x14	0.4959	0.647	0.767	0.443	-0.772	1.764
x15	-6.5932	0.705	-9.356	0.000	-7.974	-5.212
x16	2.9501	1.114	2.648	0.008	0.766	5.134
x17	-6.1937	0.623	-9.941	0.000	-7.415	-4.973
x18	-0.0587	0.587	-0.100	0.920	-1.208	1.091
x19	-10.6502	1.165	-9.143	0.000	-12.933	-8.367
x20	-4.9092	0.577	-8.515	0.000	-6.039	-3.779
x21	-7.2659	1.052	-6.909	0.000	-9.327	-5.204
x22	-6.6368	0.591	-11.221	0.000	-7.796	-5.477
x23	-0.8042	0.715	-1.125	0.260	-2.205	0.597
x24	-2.9664	0.603	-4.923	0.000	-4.147	-1.785
x25	-1.4587	0.899	-1.623	0.105	-3.220	0.302
x26	-7.2479	0.779	-9.301	0.000	-8.775	-5.721
x27	-6.2271	0.611	-10.191	0.000	-7.425	-5.029
x28	1.3294	0.259	5.124	0.000	0.821	1.838
x29	-0.5995	0.208	-2.888	0.004	-1.006	-0.193
x30	1.6216	0.259	6.254	0.000	1.113	2.130
x31	2.2811	0.279	8.175	0.000	1.734	2.828
x32	4.1627	0.340	12.258	0.000	3.497	4.828
x33	5.4057	0.401	13.464	0.000	4.619	6.193
x34	-0.5791	0.231	-2.511	0.012	-1.031	-0.127
x35	-1.3697	0.237	-5.780	0.000	-1.834	-0.905
x36	-2.1075	0.245	-8.586	0.000	-2.589	-1.626
x37	0.5394	0.270	1.994	0.046	0.009	1.069
x38	1.2180	0.340	3.588	0.000	0.553	1.883
x39	1.3501	0.402	3.361	0.001	0.563	2.137
x40	0.0662	4.099	0.016	0.987	-7.969	8.101

```

=====
Omnibus:                    58511.321    Durbin-Watson:                2.010
Prob(Omnibus):              0.000    Jarque-Bera (JB):            1819866965.133
Skew:                      19.702    Prob(JB):                    0.00
Kurtosis:                  1287.472    Cond. No.                    92.0
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly

specified.

2.1 Accessing model

```
[ ]: shap.initjs()
```

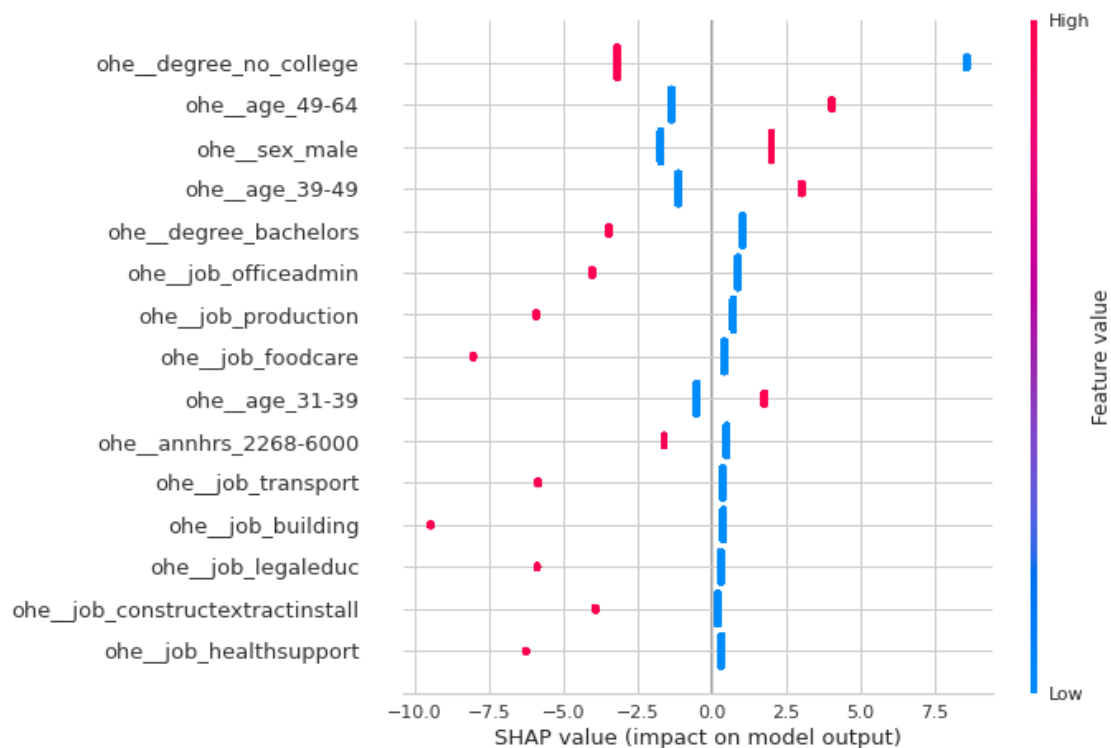
<IPython.core.display.HTML object>

```
[ ]: X_train_df = pd.DataFrame(X_train, columns=pipeline.get_feature_names_out())
X_test_df = pd.DataFrame(X_test, columns=pipeline.get_feature_names_out())
```


```
[ ]: masker = shap.maskers.Independent(data = X_train_df)
explainer = shap.LinearExplainer(lr,
    ↪masker, feature_perturbation="correlation_dependent")
shap_values = explainer.shap_values(X_test_df)
```

The `feature_perturbation` option is now deprecated in favor of using the appropriate masker (`maskers.Independent`, or `maskers.Impute`)

```
[ ]: shap.summary_plot(shap_values, X_test_df, max_display=15)
```



```
[ ]: shap.dependence_plot("ohe_sex_male", shap_values, X_test_df, x_jitter=1,
    ↪dot_size=1, interaction_index="ohe_sex_male")
```




```
[ ]: female = np.array([True if X_test[i,0] == 0 else False for i in range(X_test.
    ↳shape[0])])
```

```
[ ]: female_importances = shap_values[female].mean(axis=0)
feature_importances = pd.DataFrame(female_importances, columns =_
    ↳['mean_shap_value'])
feature_importances['feature'] = feature_names
feature_importances.sort_values('mean_shap_value', ascending=False, inplace=True)
```


```
[ ]: feature_importances[0:5].
    ↳plot(kind='barh', x='feature', y='mean_shap_value', figsize=(12,6))
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f99c98e66d0>
```

```
[ ]: feature_importances[-6:-1].
      ↪ plot(kind='barh',x='feature',y='mean_shap_value',figsize=(12,6))
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f99c35e0a00>
```



```
[ ]:
```